

Aero2Astro

Panneerselvam N

(21/08/2021)

I had trained Faster Rcnm model upto 26000 iterations. And also evaluated with both train and test dataset.

Model -1 (PP-yolo v1)

ins = 1000

bs = 8

lr = 0.0002

Training Data:

```
Accumulating evaluation results...
DONE (t=0.85s).
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.454
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.751
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.499
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.359
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.588
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.579
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.173
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.479
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.580
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.498
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.699
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.710
[08/15 07:11:53] pldet.metrics.coco_utils INFO: Per-category of bbox AP:
+-----+-----+-----+-----+-----+
| category | AP | category | AP | category | AP |
+-----+-----+-----+-----+-----+
| bird | 0.451 | book | 0.315 | bottle | 0.601 |
| car | 0.173 | chair | 0.616 | person | 0.567 |
+-----+-----+-----+-----+-----+
[08/15 07:11:53] pldet.metrics.coco_utils INFO: per-category PR curve has output to bbox
[08/15 07:11:53] pldet.engine INFO: Total sample number: 1093, average FPS: 18.4354342262
```

Test Data:

```
Accumulating evaluation results...
DONE (t=0.15s).
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.056
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.113
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.049
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.044
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.066
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.130
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.051
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.134
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.179
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.144
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.231
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.201
[08/13 14:25:45] pldet.metrics.coco_utils INFO: Per-category of bbox AP:
+-----+-----+-----+-----+-----+
| category | AP | category | AP | category | AP |
+-----+-----+-----+-----+-----+
| bird | 0.146 | book | 0.007 | bottle | 0.175 |
| car | 0.007 | person | 0.000 | chair | 0.000 |
+-----+-----+-----+-----+-----+
[08/13 14:25:45] pldet.metrics.coco_utils INFO: per-category PR curve has output to bbox_pr_curve folder.
[08/13 14:25:45] pldet.engine INFO: Total sample number: 177, average FPS: 18.61171957113853
```

Model – 2 :(PP-Yolo v2)

ins = 500

bs = 8

lr = 0.0002

Training Data :

```
Accumulating evaluation results...
DONE (t=0.64s).
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.622
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.944
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.730
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.538
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.724
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.738
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.260
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.670
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.739
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.662
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.820
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.848
[08/17 12:37:24] ppdet.metrics.coco_utils INFO: Per-category of bbox AP:
+-----+-----+-----+-----+-----+
| category | AP | category | AP | category | AP |
+-----+-----+-----+-----+-----+
| aeroplane | 0.704 | car | 0.584 | chair | 0.644 |
| cow | 0.676 | person | 0.601 | traffic_light | 0.522 |
+-----+-----+-----+-----+-----+
[08/17 12:37:24] ppdet.metrics.coco_utils INFO: per-category PR curve has output to bbox_pr_curve
[08/17 12:37:24] ppdet.engine INFO: Total sample number: 689, average FPS: 11.880296439736922
```

Test data :

```
Accumulating evaluation results...
DONE (t=0.20s).
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.191
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.362
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.167
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.117
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.257
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.307
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.137
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.334
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.409
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.259
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.490
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.643
[08/17 12:34:19] ppdet.metrics.coco_utils INFO: Per-category of bbox AP:
+-----+-----+-----+-----+-----+
| category | AP | category | AP | category | AP |
+-----+-----+-----+-----+-----+
| aeroplane | 0.343 | car | 0.100 | chair | 0.132 |
| cow | 0.352 | person | 0.055 | traffic_light | 0.164 |
+-----+-----+-----+-----+-----+
[08/17 12:34:19] ppdet.metrics.coco_utils INFO: per-category PR curve has output to bbox_pr_curve
[08/17 12:34:19] ppdet.engine INFO: Total sample number: 190, average FPS: 5.877382970685596
```


Current Model:(Faster Rcnnc)

Backbone = Faster_rcnn_X_101_32x8d_FPN_3x

bs = 4

lr = 0.001

ins = 500

Previous Checkpoints:21399 itr

Threshold = 0.85

Test Data:

```
[08/20 10:59:48 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.02 seconds
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.392
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.573
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.426
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.245
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.517
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.582
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.224
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.553
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.600
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.437
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.710
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.831
[08/20 10:59:48 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|-----|-----|-----|-----|-----|-----|
| 39.233 | 57.261 | 42.562 | 24.454 | 51.732 | 58.184 |
[08/20 10:59:48 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|-----|-----|-----|-----|-----|-----|
| aeroplane | 67.082 | car | 26.753 | chair | 31.917 |
| cow | 57.218 | person | 22.428 | traffic_light | 30.000 |
OrderedDict([('bbox',
```

```
[08/20 10:59:47 d2.evaluation.evaluator]: Total inference time: 0:00:48.845520 (0.264030 s / iter per device, on 1 devices)
[08/20 10:59:47 d2.evaluation.evaluator]: Total inference pure compute time: 0:00:48 (0.261345 s / iter per device, on 1 devices)
```

Train Data:

```
[08/20 11:06:24 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.02 seconds
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.592
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.801
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.659
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.495
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.686
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.773
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.299
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.684
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.728
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.635
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.801
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.884
[08/20 11:06:24 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|-----|-----|-----|-----|-----|-----|
| 59.158 | 80.120 | 65.930 | 49.466 | 68.642 | 77.273 |
[08/20 11:06:24 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|-----|-----|-----|-----|-----|-----|
| aeroplane | 85.382 | car | 50.095 | chair | 61.520 |
| cow | 70.940 | person | 43.778 | traffic_light | 43.236 |
OrderedDict([('bbox',
```

```
[08/20 11:06:23 d2.evaluation.evaluator]: Inference done 688/689. Data loading: 0.0018 s/iter. Inference: 0.2702 s/iter. Eval: 0.0001 s/iter.
[08/20 11:06:24 d2.evaluation.evaluator]: Total inference time: 0:03:06.400094 (0.272515 s / iter per device, on 1 devices)
[08/20 11:06:24 d2.evaluation.evaluator]: Total inference pure compute time: 0:03:04 (0.270219 s / iter per device, on 1 devices)
```


Threshold = 0.65

Test :

```
[08/21 05:09:30 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.02 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.393
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.573
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.426
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.245
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.518
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.581
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.224
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.554
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.601
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.438
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.711
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.830
[08/21 05:09:30 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|-----|-----|-----|-----|-----|-----|
| 39.256 | 57.253 | 42.563 | 24.497 | 51.781 | 58.097 |
[08/21 05:09:30 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|-----|-----|-----|-----|-----|-----|
| aeroplane | 66.918 | car | 26.636 | chair | 31.921 |
| cow | 57.237 | person | 22.582 | traffic_light | 30.240 |
OrderedDict([('bbox',
```

```
: Total inference time: 0:00:50.954686 (0.275431 s / iter per device, on 1 devices)
: Total inference pure compute time: 0:00:50 (0.272801 s / iter per device, on 1 devices)
```

Train :

```
[08/21 05:13:57 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.05 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.592
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.801
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.659
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.495
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.687
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.299
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.684
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.729
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.635
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.801
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.886
[08/21 05:13:57 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|-----|-----|-----|-----|-----|-----|
| 59.178 | 80.127 | 65.924 | 49.469 | 68.711 | 77.495 |
[08/21 05:13:57 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|-----|-----|-----|-----|-----|-----|
| aeroplane | 85.417 | car | 50.133 | chair | 61.528 |
| cow | 70.849 | person | 43.816 | traffic_light | 43.322 |
OrderedDict([('bbox',
```

```
Total inference time: 0:03:08.707873 (0.275889 s / iter per device, on 1 devices)
Total inference pure compute time: 0:03:06 (0.273226 s / iter per device, on 1 devices)
```

Current Checkpoint : 26000 its

Threshold : 0.85

Test :

```
[08/22 03:29:53 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.02
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.393
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.572
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.427
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.245
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.515
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.591
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.225
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.554
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.601
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.437
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.708
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.840
[08/22 03:29:53 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----|:-----|:-----|:-----|:-----|:-----|
| 39.268 | 57.179 | 42.705 | 24.506 | 51.460 | 59.133 |
[08/22 03:29:53 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|:-----|:-----|:-----|:-----|:-----|:-----|
| aeroplane | 67.237 | car | 26.394 | chair | 31.870 |
| cow | 57.215 | person | 22.758 | traffic_light | 30.134 |
OrderedDict([('bbox',
```

```
Total inference time: 0:00:50.160103 (0.271136 s / iter per device, on 1 devices)
Total inference pure compute time: 0:00:49 (0.268279 s / iter per device, on 1 devices)
Only: Preparing results for COCO format
```

Train :

```
[08/22 03:47:48 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[08/22 03:47:48 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.04 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.597
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.807
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.663
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.499
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.696
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.778
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.299
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.687
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.731
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.636
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.806
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.888
[08/22 03:47:48 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----|:-----|:-----|:-----|:-----|:-----|
| 59.714 | 80.673 | 66.265 | 49.898 | 69.599 | 77.833 |
[08/22 03:47:48 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|:-----|:-----|:-----|:-----|:-----|:-----|
| aeroplane | 86.043 | car | 50.530 | chair | 61.988 |
| cow | 71.296 | person | 44.628 | traffic_light | 43.798 |
OrderedDict([('bbox',
```

```
Total inference time: 0:03:03.624267 (0.268457 s / iter per device, on 1 devices)
Total inference pure compute time: 0:03:02 (0.266105 s / iter per device, on 1 devices)
Only: Preparing results for COCO format
```


Threshold : 0.65

Test :

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.393
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.572
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.427
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.245
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.515
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.591
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.225
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.554
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.601
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.437
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.708
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.840
[08/22 03:51:12 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
| :-----: | :-----: | :-----: | :-----: | :-----: | :-----: |
| 39.268 | 57.179 | 42.705 | 24.506 | 51.460 | 59.133 |
[08/22 03:51:12 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
| :-----: | :-----: | :-----: | :-----: | :-----: | :-----: |
| aeroplane | 67.237 | car | 26.394 | chair | 31.870 |
| cow | 57.215 | person | 22.758 | traffic_light | 30.134 |
OrderedDict([('bbox',
```

Train :

```
[08/22 03:55:53 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.04 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.597
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.807
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.663
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.499
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.696
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.778
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.299
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.687
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.731
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.636
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.806
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.888
[08/22 03:55:53 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
| :-----: | :-----: | :-----: | :-----: | :-----: | :-----: |
| 59.714 | 80.673 | 66.265 | 49.898 | 69.599 | 77.833 |
[08/22 03:55:53 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
| :-----: | :-----: | :-----: | :-----: | :-----: | :-----: |
| aeroplane | 86.043 | car | 50.530 | chair | 61.988 |
| cow | 71.296 | person | 44.628 | traffic_light | 43.798 |
OrderedDict([('bbox',
```