

Aero2Astro

Panneerselvam N

(04/09/2021)

I had trained Faster Rcn model upto 1010000 iterations. And also evaluated with both train and test dataset.

Model:(Faster Rcn)

Backbone = Faster_rcnn_X_101_32x8d_FPN_3x

bs = 4

lr = 0.001

ins = 500

Previous Checkpoint : 83500

Threshold : 0.85

Test :

```
[09/01 11:00:32 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.05 second
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.388
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.563
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.416
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.234
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.506
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.555
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.222
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.547
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.590
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.417
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.698
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.834
[09/01 11:00:32 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|-----|-----|-----|-----|-----|-----|
| 38.774 | 56.323 | 41.635 | 23.398 | 50.569 | 55.524 |
[09/01 11:00:32 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|-----|-----|-----|-----|-----|-----|
| aeroplane | 66.686 | car | 25.895 | chair | 31.846 |
| cow | 57.283 | person | 20.700 | traffic light | 30.234 |

Total inference time: 0:02:15.337548 (0.731554 s / iter per device, on 1 devices)
Total inference pure compute time: 0:02:14 (0.728070 s / iter per device, on 1 devices)
```

Train :

```
[09/01 11:13:20 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.05 second
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.641
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.845
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.709
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.548
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.737
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.822
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.309
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.714
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.755
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.665
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.827
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.915
[09/01 11:13:20 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|-----|-----|-----|-----|-----|-----|
| 64.101 | 84.488 | 70.862 | 54.822 | 73.704 | 82.153 |
[09/01 11:13:20 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|-----|-----|-----|-----|-----|-----|
| aeroplane | 88.188 | car | 55.989 | chair | 67.044 |
| cow | 73.846 | person | 51.198 | traffic_light | 48.340 |
```

```
Total inference time: 0:08:20.223715 (0.731321 s / iter per device, on 1 devices)
Total inference pure compute time: 0:08:18 (0.728191 s / iter per device, on 1 devices)
tionl: Preparing results for COCO format
```

Threshold : 0.65

Test :

```
[09/01 11:24:43 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[09/01 11:24:43 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.02 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.388
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.563
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.416
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.234
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.506
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.555
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.222
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.547
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.590
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.417
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.698
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.834
[09/01 11:24:43 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 38.774 | 56.323 | 41.635 | 23.398 | 50.569 | 55.524 |
[09/01 11:24:43 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| aeroplane | 66.686 | car | 25.895 | chair | 31.846 |
| cow | 57.283 | person | 20.700 | traffic_light | 30.234 |
OrderedDict([('bbox',
```

```
: Total inference time: 0:02:16.642567 (0.738608 s / iter per device, on 1 devices)
: Total inference pure compute time: 0:02:15 (0.735123 s / iter per device, on 1 devices)
```

Train :

```
[09/01 11:35:26 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.05 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.641
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.845
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.709
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.548
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.737
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.822
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.309
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.714
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.755
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.665
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.827
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.915
[09/01 11:35:26 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 64.101 | 84.488 | 70.862 | 54.822 | 73.704 | 82.153 |
[09/01 11:35:26 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| aeroplane | 88.188 | car | 55.989 | chair | 67.044 |
| cow | 73.846 | person | 51.198 | traffic_light | 48.340 |
OrderedDict([('bbox',
```

```
Total inference time: 0:08:12.162782 (0.719536 s / iter per device, on 1 devices)
Total inference pure compute time: 0:08:10 (0.716395 s / iter per device, on 1 devices)
tionl: Preparing results for COCO format
```


Previous checkpoints : 100999 itr (with Horizontal flip augmentation)

Threshold : 0.85

Test :

```
[09/03 13:06:29 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.02 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.388
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.575
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.420
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.238
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.503
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.589
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.221
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.546
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.589
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.428
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.692
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.827
[09/03 13:06:29 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|-----|-----|-----|-----|-----|-----|
| 38.795 | 57.503 | 42.001 | 23.815 | 50.281 | 58.874 |
[09/03 13:06:29 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|-----|-----|-----|-----|-----|-----|
| aeroplane | 66.272 | car | 27.213 | chair | 31.770 |
| cow | 57.276 | person | 20.187 | traffic_light | 30.051 |
OrderedDict([('bbox',
```

```
Total inference time: 0:02:27.579475 (0.797727 s / iter per device, on 1 devices)
Total inference pure compute time: 0:02:26 (0.794334 s / iter per device, on 1 devices)
only: Preparing results for COCO format
```

Train :

```
[09/03 13:17:54 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.04 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.643
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.851
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.719
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.553
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.734
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.814
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.305
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.713
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.753
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.666
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.823
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.906
[09/03 13:17:54 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|-----|-----|-----|-----|-----|-----|
| 64.338 | 85.093 | 71.879 | 55.269 | 73.435 | 81.437 |
[09/03 13:17:54 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|-----|-----|-----|-----|-----|-----|
| aeroplane | 87.213 | car | 56.145 | chair | 67.627 |
| cow | 73.221 | person | 52.803 | traffic_light | 49.017 |
```

```
Total inference time: 0:09:01.542198 (0.791728 s / iter per device, on 1 devices)
Total inference pure compute time: 0:08:59 (0.788535 s / iter per device, on 1 devices)
```

Threshold : 0.65

Test:

```
[09/03 13:20:57 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.02 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.388
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.575
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.420
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.238
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.503
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.589
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.221
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.546
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.589
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.428
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.692
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.827
[09/03 13:20:57 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|-----|-----|-----|-----|-----|-----|
| 38.795 | 57.503 | 42.001 | 23.815 | 50.281 | 58.874 |
[09/03 13:20:57 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|-----|-----|-----|-----|-----|-----|
| aeroplane | 66.272 | car | 27.213 | chair | 31.770 |
| cow | 57.276 | person | 20.187 | traffic light | 30.051 |
```

```
Total inference time: 0:02:27.451421 (0.797035 s / iter per device, on 1 devices)
Total inference pure compute time: 0:02:26 (0.793703 s / iter per device, on 1 devices)
```

Train :

```
[09/03 13:30:22 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.02 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.643
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.851
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.719
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.553
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.734
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.814
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.305
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.713
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.753
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.666
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.823
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.906
[09/03 13:30:22 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|-----|-----|-----|-----|-----|-----|
| 64.338 | 85.093 | 71.879 | 55.269 | 73.435 | 81.437 |
[09/03 13:30:22 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|-----|-----|-----|-----|-----|-----|
| aeroplane | 87.213 | car | 56.145 | chair | 67.627 |
| cow | 73.221 | person | 52.803 | traffic_light | 49.017 |
```

```
Total inference time: 0:09:01.864907 (0.792200 s / iter per device, on 1 devices)
Total inference pure compute time: 0:08:59 (0.788929 s / iter per device, on 1 devices)
```


Last checkpoints: 109999

Threshold : 0.85

Test:

```
[09/05 01:03:45 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.02 seconds
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.387
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.573
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.421
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.239
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.505
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.588
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.222
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.547
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.591
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.432
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.693
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.824
[09/05 01:03:45 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 38.740 | 57.349 | 42.095 | 23.868 | 50.457 | 58.834 |
[09/05 01:03:45 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| aeroplane | 65.832 | car | 27.135 | chair | 31.566 |
| cow | 57.669 | person | 20.150 | traffic_light | 30.084 |
```

```
Total inference time: 0:02:26.242037 (0.790497 s / iter per device, on 1 devices)
Total inference pure compute time: 0:02:25 (0.787103 s / iter per device, on 1 devices)
```

Train :

```
[09/05 01:13:33 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.04 seconds
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.645
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.856
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.716
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.558
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.734
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.817
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.306
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.713
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.754
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.669
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.823
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.909
[09/05 01:13:33 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 64.504 | 85.629 | 71.646 | 55.846 | 73.401 | 81.655 |
[09/05 01:13:33 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| aeroplane | 86.627 | car | 56.614 | chair | 67.623 |
| cow | 73.512 | person | 53.348 | traffic_light | 49.303 |
```

```
Total inference time: 0:08:57.177787 (0.785348 s / iter per device, on 1 devices)
Total inference pure compute time: 0:08:55 (0.782235 s / iter per device, on 1 devices)
```

Threshold: 0.65

Test:

```
[09/05 01:17:41 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.02 second
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.387
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.573
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.421
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.239
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.505
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.588
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.222
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.547
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.591
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.432
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.693
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.824
[09/05 01:17:41 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
| :-----: | :-----: | :-----: | :-----: | :-----: | :-----: |
| 38.740 | 57.349 | 42.095 | 23.868 | 50.457 | 58.834 |
[09/05 01:17:41 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
| :-----: | :-----: | :-----: | :-----: | :-----: | :-----: |
| aeroplane | 65.832 | car | 27.135 | chair | 31.566 |
| cow | 57.669 | person | 20.150 | traffic_light | 30.084 |
OrderedDict([('bbox',
```

```
Total inference time: 0:02:26.359061 (0.791130 s / iter per device, on 1 devices)
Total inference pure compute time: 0:02:25 (0.787750 s / iter per device, on 1 devices)
```

Train :

```
[09/05 01:28:46 d2.evaluation.fast_eval_api]: COCOeval_opt.accumulate() finished in 0.04 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.645
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.856
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.716
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.558
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.734
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.817
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.306
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.713
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.754
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.669
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.823
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.909
[09/05 01:28:46 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
| :-----: | :-----: | :-----: | :-----: | :-----: | :-----: |
| 64.504 | 85.629 | 71.646 | 55.846 | 73.401 | 81.655 |
[09/05 01:28:46 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
| :-----: | :-----: | :-----: | :-----: | :-----: | :-----: |
| aeroplane | 86.627 | car | 56.614 | chair | 67.623 |
| cow | 73.512 | person | 53.348 | traffic_light | 49.303 |
OrderedDict([('bbox',
```

```
Total inference time: 0:08:57.562220 (0.785910 s / iter per device, on 1 devices)
Total inference pure compute time: 0:08:55 (0.782757 s / iter per device, on 1 devices)
```