

Panneerselvam N
(09/09/2021)

Today i had trained faster rcnn model upto 149499 iterations with more augmementation.

Previous Checkpoint: 138000

Test :

Augmentation : Resize(1024), HFlip

```
[09/09 03:58:44 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.02 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.384
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.565
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.415
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.232
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.503
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.552
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.221
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.541
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.581
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.404
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.688
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.819
[09/09 03:58:44 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|-----|-----|-----|-----|-----|-----|
| 38.352 | 56.507 | 41.474 | 23.209 | 50.253 | 55.211 |
[09/09 03:58:44 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|-----|-----|-----|-----|-----|-----|
| aeroplane | 64.450 | car | 26.561 | chair | 30.293 |
| cow | 58.988 | person | 20.008 | traffic_light | 29.813 |
```

```
Total inference time: 0:02:17.554629 (0.743539 s / iter per device, on 1 devices)
Total inference pure compute time: 0:02:16 (0.740199 s / iter per device, on 1 devices)
```

Train :

```
[09/09 04:08:01 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.05 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.656
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.865
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.733
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.565
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.743
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.825
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.306
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.717
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.756
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.667
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.827
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.912
[09/09 04:08:01 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|-----|-----|-----|-----|-----|-----|
| 65.610 | 86.528 | 73.318 | 56.502 | 74.312 | 82.524 |
[09/09 04:08:01 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|-----|-----|-----|-----|-----|-----|
| aeroplane | 88.144 | car | 57.644 | chair | 67.827 |
| cow | 74.917 | person | 55.569 | traffic_light | 49.560 |
OrderedDict([('bbox',
```

```
Total inference time: 0:08:27.085180 (0.741353 s / iter per device, on 1 devices)
Total inference pure compute time: 0:08:24 (0.738200 s / iter per device, on 1 devices)
```


Current Checkpoints : 149500 itr

Augmentation : Resize(1024),Hflip ,Random brightness , Random Rotation , Random crop

Test :

```
[09/09 14:34:01 d2.evaluation.fast_eval_api]: Evaluate annotation type *bbox*
[09/09 14:34:02 d2.evaluation.fast_eval_api]: COCOeval opt.evaluate() finished in 0.11 seconds.
[09/09 14:34:02 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[09/09 14:34:02 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.02 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.203
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.520
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.086
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.156
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.287
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.287
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.127
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.338
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.374
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.295
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.426
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.498
[09/09 14:34:02 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 20.263 | 52.047 | 8.584 | 15.602 | 28.665 | 28.670 |
[09/09 14:34:02 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| aeroplane | 27.905 | car | 16.978 | chair | 16.489 |
| cow | 23.104 | person | 13.267 | traffic_light | 23.835 |
OrderedDict([('bbox',
```

```
Total inference time: 0:02:26.119563 (0.789835 s / iter per device, on 1 devices)
Total inference pure compute time: 0:02:21 (0.763777 s / iter per device, on 1 devices)
```

Train :

```
[09/09 15:03:50 d2.evaluation.fast_eval_api]: Accumulating evaluation results...
[09/09 15:03:50 d2.evaluation.fast_eval_api]: COCOeval opt.accumulate() finished in 0.07 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.321
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.801
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.116
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.311
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.375
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.392
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.163
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.413
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.447
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.412
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.478
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.518
[09/09 15:03:50 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP | AP50 | AP75 | APs | APm | APl |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 32.063 | 80.082 | 11.644 | 31.081 | 37.548 | 39.221 |
[09/09 15:03:50 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP | category | AP | category | AP |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| aeroplane | 31.379 | car | 30.986 | chair | 36.573 |
| cow | 35.506 | person | 23.276 | traffic_light | 34.657 |
OrderedDict([('bbox',
```

```
Inference done 686/689. Dataloading: 0.0023 s/iter. Inference: 0.7593 s/iter. Eval: 0.0003
Total inference time: 0:08:41.578473 (0.762542 s / iter per device, on 1 devices)
Total inference pure compute time: 0:08:39 (0.759366 s / iter per device, on 1 devices)
```